

Logbook — Anomalous transport in soaring flights

Matteo Di Sante

Working log

Working notebook. The timeline is regenerated from the git history; the narrative notes are written by hand.

1 Timeline (auto-generated from git)

Date	Event (commit subject)
2026-08-05	docs: the proposed structure for Chapter 3
2026-08-05	feat(kinematics): the last three catalogued observables, measured
2026-08-05	fix(thesis): the edge-sample effect was typed, and both figures were wrong
2026-08-05	feat(macros): the archive measurements the body still typed by hand
2026-08-05	docs: the report said the propagator was the measurement most needed next; it is built
2026-08-05	docs(logbook): the propagator is built, and what it says
2026-08-05	feat(propagator): the exponent from the bulk, on the whole archive, and what it costs the exponent
2026-08-05	data(shape): the run statistics, from an estimator that checks the runs it emits
2026-08-05	feat(propagator): an exponent from the bulk of the distribution, with neither a moment nor an origin
2026-08-05	docs: the last three audit defects, all fixed
2026-08-05	fix(audit): the departure time was read at the grid edge, and one ratio averaged two populations
2026-08-04	docs: the closing checks, and the one thing still outstanding
2026-08-04	docs: the findings table closes at thirty-three
2026-08-04	docs: the correctness audit and the invalid null, in the report
2026-08-04	fix(synthetic,bootstrap,shape): three latent defects, two of them in the validation itself
2026-08-04	fix(audit): the runs test was read against a null the data cannot satisfy
2026-08-04	fix(moments): the departure from a line was a standard deviation called an rms
2026-08-04	fix(persistence): the run scan validated only the runs it happened to extend
2026-08-04	docs(logbook): the non-Gaussian result, with the two disciplines kept apart
2026-08-04	docs: the findings table catches up with the last round
2026-08-04	docs(regenerate): state what the two shape steps actually cost
2026-08-04	fix(thesis): two different quantities were both called alpha_2, seventeen lines apart
2026-08-04	docs: the fourth self-correction, on a paragraph this pass wrote
2026-08-04	docs(thesis): four of the six implementation pointers in Chapter 3 landed on nothing

Date	Event (commit subject)
2026-08-04	fix(thesis): the non-Gaussian calibration was mine, not the test's, and it flattened the two disciplines
2026-08-04	docs: section 3.2 is no longer the gap in the pass
2026-08-04	fix(thesis): ten quantitative slips, most of them a macro away from being right
2026-08-04	fix(thesis): fifteen more claims the document contradicts, from a full sentence-level review
2026-08-04	docs: fill in the closing checks of the verification report
2026-08-04	docs(thesis): cite the FAI Sporting Code, and drop two sentences that restate their neighbours
2026-08-04	fix(wavelet): the logscale diagram had the wrong convention and a low-biased fit
2026-08-04	fix(thesis): say what the acquisition check actually checks, and six other overstatements
2026-08-04	fix(thesis): the front matter described a document that no longer exists
2026-08-04	feat(shape): the paraglider pass lands, and the thesis builds again
2026-08-04	docs: record that verify_dataset was started and stopped
2026-08-04	docs: the docs-build findings in the verification report
2026-08-04	fix(docs): make 'mkdocs build --strict' pass, and annotate the analysis signatures
2026-08-04	docs: three more findings in the verification report
2026-08-04	fix(thesis): the shape caption described four panels; the figure has six
2026-08-04	style(thesis): ragged-right the third column of the tooling table
2026-08-04	docs(thesis): state the agreement with the published H at the precision it is quoted to
2026-08-04	feat(transport): compare the exponent against the published value, through macros
2026-08-04	docs: the verification report
2026-08-04	fix(thesis): the abstract described a plan the document no longer is; repair a \rev group
2026-08-04	docs(logbook): first-passage times tried and rejected, with the reason
2026-08-04	docs(thesis): record why the first-passage exponent is not an independent check
2026-08-04	fix(thesis): the introduction pointed the modelling decision at the wrong chapter
2026-08-04	fix(thesis): the exponent figure was generated and never shown; split three page-long paragraphs
2026-08-04	docs(guide): refresh the schema dumps, which were a pipeline run behind
2026-08-04	docs(logbook): bring Next steps back in step with Chapter 3, and fix a dead path
2026-08-04	docs(thesis): withdraw the Gaussian reading, and answer the memory-versus-exponent question
2026-08-04	feat(check): find measured numbers typed where a macro already carries them
2026-08-04	feat(shape): measure the non-Gaussian parameter, and the tail of the memory
2026-08-04	feat(observables): replace the unused Green-Kubo integral with the bound the data supports
2026-08-04	perf(persistence): vectorise the run scan, and prove it is the same scan
2026-08-04	chore: close the gaps the verification pass found in the code

Date	Event (commit subject)
2026-08-04	feat(observables): the moment spectrum on the archive says it is not a Levy walk
2026-08-04	feat(observables): velocity memory, persistence runs, and the bias that costs a whole unit
2026-08-04	feat(observables): the moment spectrum, and the filter order it forces a choice about
2026-08-04	thesis: set the narrow table columns ragged right
2026-08-04	fix(hooks): the drift check named a file the chapter split removed
2026-08-04	docs: bring the reference, the chain and the logbook up to the code
2026-08-04	thesis(ch3): write the measurement the chapter was missing
2026-08-04	feat(transport): stratify the exponent, and calibrate what $\alpha_2 = 2$ means
2026-08-04	fix(transport): spell out the filter order, and evict five phantom duplicates
2026-08-04	feat(transport): drop the "clean range" idea, which the data does not support
2026-08-04	feat(transport): reduce the variation pass, and find that the sampling error is not the error
2026-08-04	feat(stats): resample clusters, not flights, and supply a null the shape cannot leak into
2026-08-04	feat(observables): measure filtered variations per flight, and find the task geometry
2026-08-04	feat(observables): fit the number of regimes, and refuse to trust a circular null
2026-08-04	feat(observables): validated synthetic processes, and the case against centring
2026-08-04	thesis: split the planning chapter into global transport and flight phases
2026-08-04	docs(plan): add the Chapter 3 inventory, which overrules the earlier sketches
2026-08-04	fix(audit): test the arrangement of the residuals, not only their size
2026-08-04	feat(dataset): read the discards as a cascade, and check them against the calendar
2026-08-04	refactor(analysis): separate computing from drawing, split the grab-bag module
2026-08-04	chore: drop what the repo no longer needs; break the oversized parameter table
2026-08-04	chore: regenerate the committed artefacts
2026-08-04	docs(logbook): record the pipeline build and the MSD audit, realign Next steps
2026-08-04	thesis(ch3): report the two MSD estimators as two measurements, not one
2026-08-04	thesis(ch2): the pipeline as built, the analysis ensemble, and Sec. 2.8 in full
2026-08-04	docs: the on-disk layout, the regeneration order and the handoff to the MSD
2026-08-04	feat(prelim): the preliminary characterization of Sec. 2.8, on real maps
2026-08-04	feat(audit): keep what the MSD averages away, and use it to test the curve
2026-08-04	feat(reporting): the pipeline census, the macro contract and the regeneration chain
2026-08-04	feat(verify): check the written tables against the invariants, and against the code
2026-08-04	feat(transport): the two mean-squared-displacement estimators
2026-08-04	feat(analysis): stream whole flights out of the fix table
2026-08-04	feat(preproc): the seven-stage trajectory pipeline
2026-08-04	feat(igc): tolerate partial altitude fields and the midnight roll-over

Date	Event (commit subject)
2026-08-04	chore(tooling): require Python 3.12, fix the mypy roots, sweep the formatting
2026-08-03	chore: regenerate the committed artefacts
2026-08-03	docs: align the pipeline guide, complete the bibliography inventory
2026-08-03	docs(thesis): apply the second review pass
2026-08-03	feat(census): count the placeholder dates, per season
2026-08-03	feat(preproc): altitude-activity cut 75 -> 600 m
2026-08-03	fix(analysis): mark the gap cut at the bound actually applied
2026-08-03	feat(analysis): draw the third flight-level diagnostic
2026-08-03	fix(analysis): judge barometric presence strictly, from one definition
2026-07-24	docs(thesis): review-pass revisions, geodesy appendix and pipeline map
2026-07-24	docs(preproc): align pipeline guide with 20 km, \Preproc macros and the contract
2026-07-24	feat(preproc): 20 km path cut; export thresholds and census as macros
2026-07-13	docs(thesis): trim next-steps plan; sync logbook
2026-07-13	docs(thesis): running headers and typographic refinements
2026-07-13	docs(thesis): voice and flow pass; fix EB limit in CTRW appendix
2026-07-13	docs(thesis): revise dataset chapter after review notes
2026-07-09	test: pin p-dependence with invariants instead of a derived constant
2026-07-09	test: strengthen conservativeness check; cover the ~ expansion branch
2026-07-09	test: cover altitude-noise statistics/PSD aggregation and load_config
2026-07-09	test: decouple suite from editable .pth; fix stale Config() calls
2026-07-09	refactor(config): remove paraglider-first bias from load_config
2026-07-07	logbook: seventh audit entry (global sweep)
2026-07-07	audit(global): consistency sweep across every document
2026-07-07	logbook: sixth audit entry
2026-07-07	preproc(final): every hyperparameter pinned; tau_c measured; flight filter corrected on census evidence
2026-07-07	logbook: fifth audit entry; unbreakable identifier in C2 pointed at the repro table
2026-07-07	preproc(audit): implementer's walkthrough — four logical bugs closed; C2 de-walled
2026-07-07	logbook: fourth audit entry
2026-07-07	preproc(audit): measure the gap cap on the census; suspect-wait flag; blueprint realigned
2026-07-07	logbook: third audit entry
2026-07-07	preproc(audit): 2*dt floor on the gap cap; interpolated mask; interior-ground detector; per-segment gates
2026-07-07	logbook: second split audit entry
2026-07-07	preproc(gaps): absolute cap on bridgeable gaps; split mechanics made concrete
2026-07-07	logbook: four July entries; Next steps renumbered to the merged layout
2026-07-07	preproc(split): state what a split means for each observable; fix the segment gate
2026-07-07	style(voice): sweep appendices B and C2
2026-07-07	style(voice): sweep the body chapters
2026-07-07	style(voice): sober headings and prose in the dataset chapter
2026-07-07	impl(params): the fix-level working defaults as a table
2026-07-07	thesis(enu): signpost the two-step conversion
2026-07-07	impl(validation): the five cleaning checks as a numbered list

Date	Event (commit subject)
2026-07-07	thesis(structure): fold data acquisition into the dataset chapter; let the text breathe
2026-07-07	cleaning(frozen-lock): rank the witness; quantify the body; lighten tab:cleaning
2026-07-07	ffvl link fixed
2026-07-07	thesis(cleaning-spec): one contract across body, appendix and blueprint
2026-07-07	reporting(census): generate the census-scan numbers as \StatScan* macros
2026-07-07	config(preproc): condition the vertical savgol timescale on alt_source; fix stale refs
2026-07-07	cleaning: make fix-level cleaning conservative — never lose a search/climb fix
2026-07-07	fix(altitude-PSD): correct the baro-vs-GNSS noise claim; robust median+band figure
2026-07-07	bibliography: ground the transport-analysis & CTRW methods
2026-07-07	bibliography: ground the pre-processing methods (Savitzky–Golay, Welch, WGS84, Hampel/MAD)
2026-07-07	bibliography: verify every entry, archive sources locally, fix the FAI citation
2026-07-07	polish(prose+typography): tighten cleaning sections, clear all overfull boxes
2026-07-07	logbook: record the fix-level cleaning redesign
2026-07-07	redesign fix-level cleaning: three-detector taxonomy, clear removal semantics
2026-07-07	fix(typography): float placement, long-identifier wrapping, A/V prose
2026-07-07	restructure thesis: narrative body + implementation-details appendix
2026-07-06	fix(gap-fig): tighten Fig 4.6 panel (b) x-range; pin all figures in place
2026-07-06	fix(gap-fig): linear axis/bins for Fig 4.6 panel (a), reveals real small-integer structure
2026-07-06	fix(sampling-fig): discrete bins for Fig 4.7 (native dt), fraction not density
2026-07-06	fix(fixlevel-fig): shared 13 m/s vertical-speed bound; fix a general long-caption LaTeX limit
2026-07-06	fix(fixlevel-fig): integer-aligned bins for vz; revise horizontal-speed bounds to 45/55
2026-07-06	fix(fixlevel-fig): widen x-range for vxy/vz tails; explain vz’s non-empty valleys
2026-07-06	feat(preproc): per-discipline fix-level speed bounds, on a much larger sample
2026-07-06	fix(altitude-noise): Fig 4.1 panel (a) back to dual y-axis, over the full 30-min window
2026-07-06	fix(thesis): Fig 4.1 redesign (drift over jitter), Appendix B panel-letter bug, missing sample sizes
2026-07-06	chore: add MIT LICENSE, README badges, and a real test-running CI
2026-07-06	ci(docs): auto-deploy on push again, scoped to paths that feed the site
2026-07-06	docs(logbook): reformat narrative notes for readability; re-sync Next steps; docs deploy on demand
2026-07-06	fix(altitude-noise): dual-axis + decoupled windows for Fig. 4.1 panels (a)/(b); chore(deps): dev/analysis as uv default groups
2026-07-06	docs: data-driven fix-level section; four-panel altnoise caption; blueprint, logbook, README
2026-07-06	chore(data): explicit SSD->repo seasons_index sync; document the offline snapshot
2026-07-06	feat(altitude-noise): four-panel figure (raw channels, difference, PSD, availability)

Date	Event (commit subject)
2026-07-06	feat(preproc): data-driven fix-level diagnostics; gap-histogram readability; drop redundant time-gap
2026-07-06	fix(igc): treat only full-day drops as a midnight roll-over
2026-07-06	chore: SSD data layout (raw/catalog/derived) + pyarrow for the scan cache
2026-07-05	thesis ch.4: data-driven pre-processing pipeline + geodesy figures
2026-07-02	Symmetric para/hang dataset; adopt barometric altitude; IGC parser + noise diagnostics
2026-07-01	Rename soaring-ffvl -> soaring-para; add hang-glider data and update all docs
2026-07-01	Fix delta config: season start 2001, data_root /Volumes/SSD_DISANTE/hang_glidern/delta_cfd_igc
2026-07-01	Add soaring-delta CLI for hang-glider CFD data (delta.ffvl.fr)
2026-07-01	Add analysis sub-package and preprocessing diagnostics; revise thesis chapters 1-4
2026-07-01	Expand pre-processing: explicit cleaning and filtering sections
2026-07-01	Clarify stationarity/compatibility check paragraph
2026-07-01	Add preliminary characterization section before main observables
2026-07-01	Add isotropy test protocol to the propagator item
2026-07-01	Rewrite propagator item: formal derivation of marginal scaling
2026-07-01	Reorganize TA-MSD item for clarity and typographic structure
2026-07-01	Clarify ergodicity breaking and add CTRW appendix
2026-06-28	Revise chapter 4 observables and the IGC description
2026-06-28	Warn when thesis and logbook next-steps drift apart
2026-06-28	Align logbook next steps with the thesis planning chapter
2026-06-28	Publish the logbook alongside the thesis
2026-06-28	Add living thesis document and build automation
2026-06-26	Deploy docs via GitHub Pages Actions artifact instead of gh-pages branch
2026-06-26	Fail fast with a clear message when data_root is unset
2026-06-26	Add 'verify' subcommand and remove machine-specific data_root path
2026-06-26	Link published docs in README; drop Development section
2026-06-25	Add GitHub Pages deployment via gh-pages branch
2026-06-25	Translate all documentation, comments, and strings to English
2026-06-25	Initialize thesis repo: FFVL CFD .igc data acquisition

2 Narrative notes

Section and figure numbers inside dated entries refer to the document layout at the time of writing; the thesis has since been restructured (the dataset chapter is now Ch. 2, “Next steps” Ch. 3).

2026-06 — Paraglider data acquisition complete. The full FFVL paraglider CFD archive (`parapente.ffvl.fr`, seasons 1999–2000 to 2025–2026) has been scraped: about 203,000 flight records, of which ~186,000 carry a GPS track. All 186,052 available tracks were downloaded.

What worked Impersonating a Chrome TLS fingerprint (`curl_cffi`) reliably passes the FFVL Cloudflare challenge; the resumable, atomic, validated download survived many interruptions without corruption.

Gotchas 173 tracks could not be obtained — 172 because the server returns something that is not a valid IGC (broken or placeholder links on FFVL’s side, so a retry does not help), and

1 due to a transient Cloudflare challenge (worth a later retry). On macOS/exFAT the OS scatters `._*` sidecar files next to every written file; a `clean` step removes them.

2026-06 — Integrity verified. A full on-disk re-scan (`verify`) confirmed that all 186,052 downloaded files are structurally valid IGC: no truncated files, no leftover `.part`, no read errors. The verification logic was promoted from a throwaway script to a first-class `soaring-para verify` subcommand (reusing the download-time validator), with tests.

2026-06 — Robustness of the data path. The external disk remounted under a different name (`HDD_DISANTE` → `SSD_DISANTE`), which would have silently broken a hard-coded path. Fix: removed every machine-specific path from the repo (placeholder + `SOARING_PARA_DATA_ROOT` environment variable) and added a start-up guard that aborts with a clear, actionable message when `data_root` is unset/unmounted — and verified that `/Volumes` is not user-writable, so a wrong path can never cause a silent download to the wrong place.

2026-07 — Hang-glider data acquired; naming corrected.

Second source added The hang-glider CFD (`delta.ffvl.fr`, seasons 2001–2002 to 2025–2026), sharing the acquisition code (`soaring.acquisition.ffvl`) via a second CLI entry point `soaring-delta` (config `delta_download.yaml`, env var `SOARING_DELTA_DATA_ROOT`). Data: ~9,300 flights, 6,716 tracks downloaded, 30 failures (broken server-side links, same pattern as paragliders). Stored at `/Volumes/SSD_DISANTE/hang_gliders/delta_cfd_igc/`.

Naming corrected The paraglider CLI was renamed from the generic `soaring-ffvl` to `soaring-para` (and `SOARING_FFVL_DATA_ROOT` → `SOARING_PARA_DATA_ROOT`), and its config from `ffvl_download.yaml` to `para_download.yaml`, to avoid ambiguity now that both sources are present.

2026-07 — Hang-glider catalog built; datasets reorganized symmetrically. Ran the full pipeline on the hang-glider data (`soaring-delta clean, verify, build-catalog, status`), which so far had only been downloaded: 9,259 flights, 6,746 with a track, 6,716 downloaded and all verified (0 integrity failures), 30 unrecoverable server-side links.

Data layout The local `data/` folder now mirrors the SSD — one directory per discipline (`paragliders/`, `hang_gliders/`), each with a versioned `seasons_index.csv` and a git-ignored `catalog.csv`.

Thesis symmetry `generate_stats.py` now reads both indices and emits per-discipline macros (`\StatPara*`, `\StatHang*`) and two per-season tables; the first three thesis chapters were made symmetric across both disciplines.

2026-07 — Altitude channel decision: barometric. Decided to take the vertical dynamics from the *barometric* altitude rather than the GNSS altitude, diverging from the reference pipeline (Jérémy's code derives z and v_z from `AltGNSS`, reading but not using `AltPre`).

Rationale The pressure altitude has sub-metre relative precision and low high-frequency noise, its errors being slow (offset, drift) and killed by differencing, whereas the GNSS vertical carries a high-frequency noise floor that derivatives amplify. Confirmed empirically with a new noise diagnostic — an IGC B-record parser (`soaring.analysis.igc`) and an altitude-noise analysis (`soaring.analysis.altitude_noise`) whose Welch power spectra show the GNSS floor two to three orders of magnitude above the barometric one near Nyquist (thesis appendix on the method).

Consequences The A/V validity flag is subsumed by the missing-altitude check on the adopted channel (a V fix keeps position + baro on a baro flight, but is dropped on a GNSS-fallback flight); one altitude channel per trajectory, no baro/GNSS splicing; flights with no pressure

sensor fall back to unfiltered GNSS, tagged by a per-flight `alt_source` column; the trimming threshold is on horizontal speed only ($v_{xy} > 5$ m/s, sustained); cleaning/trimming run on the raw geographic coordinates (great-circle speeds), with the ENU conversion done *last*, origin at take-off.

2026-07 — Baro-absence fraction: sized sampling, not a full sweep. A full-population parse of the paraglider archive (186,052 files) for the baro-channel-absence figure took an unacceptable ~ 45 min serially; parallelising across 8 processes cut this to ~ 15 min, still too slow to run routinely.

Fix A properly justified *simple random sample*: the standard proportion sample-size formula ($n = z^2 p(1-p)/e^2$, conservative $p = 0.5$) gives $n = 2401$ for a target 95%-confidence margin of error of ± 2 percentage points (`required_sample_size` in `soaring.analysis.altitude_noise`, with `proportion_ci` reporting the point estimate and interval). The population is treated as effectively infinite: at $n = 2401$ out of 186,052 flights, the finite-population correction would move n by only ~ 30 , not worth the extra parameter. The hang-glider archive (6716 files) stays a full census (cheap enough, under two minutes). Total runtime: ~ 50 s. Result: paragliders $(28.4 \pm 1.8)\%$ (95% CI, sampled), hang gliders 17.5% (exact, census).

Also caught mid-flight The IGC parser accepted syntactically well-formed but semantically invalid B records (e.g. minute ≥ 60 , hour ≥ 24 , out-of-range latitude/longitude) without rejecting them; added explicit format-validity checks (9 new tests) and corrected the thesis wording, which had conflated this *implemented* format-validity layer with the fix-level *dynamics-plausibility* layer of Table 4.1 – at the time designed (a dataclass of thresholds) but not yet implemented as an enforced filter.

2026-07 — Pre-processing made data-driven and config-driven. Reworked Chapter 4 so the flight-level filtering is computed from the *full parsed dataset* (paragliders and hang gliders), not the declared catalog metadata.

What changed A track scan (`soaring.analysis.census.scan_tracks`) parses every downloaded track; Figure 4.2 shows the real duration and flown-path-length distributions with per-variable *marginal* retention curves. Adopted cuts: duration ≥ 40 min and path length ≥ 30 km (a minimal cut, $\sim 95\%$ retained); the earlier min-fix-count cut was dropped as redundant with duration, and the *extent* proxy was replaced by total flown *path length* (median 86 km para, 158 km hang).

Key decision Every pre-processing threshold now lives in `configs/preprocessing.yaml` (never hard-coded), loaded via `load_preproc_config` into typed dataclasses.

Also added Great-circle (haversine) inter-fix speed/distance *before* the ENU conversion; a Savitzky–Golay smoothing/differentiation section with a noise-matched procedure for its two hyperparameters; the switch from a common resampling cadence to native per-flight Δt (Figure 4.5); and an implementation blueprint (`docs/guide/preprocessing-pipeline.md`) recording the exact schema and steps, to retire into code once the pipeline is built.

2026-07 — Gap diagnostics; scan cache; baro-presence consolidated. Added Figure 4.5, the sampling-regularity diagnostics (largest gap relative to native Δt , missing fraction of a uniform grid, each with its marginal retention curve), fulfilling a promise Sec. 4.1.6 already made but had not yet delivered.

Scan cache A fresh full parse would have repeated the ~ 15 – 20 min paraglider census just done for Figure 4.2, so the per-flight scan is now **cached** to `data/<discipline>/track_scan.parquet` (gitignored, like `catalog.csv`; `load_or_scan_tracks` reads it if present, else scans and writes it — later moved to the SSD, see below). `track_stats` was extended with a few near-free

byproducts of the same scan — barometric-presence fraction, max horizontal/vertical speed, barometric altitude range — for future validation of Table 4.1’s fix-level bounds.

Consolidation The barometric-presence fraction (Sec. 4.1.1) now reads from this same cache (`altitude_noise.baro_presence_from_scan`) instead of running its own separate scan, turning the paraglider number from a sized sample into an exact census at no extra cost: 30.1 % (exact, $n = 186,025$) — consistent with the earlier sampled estimate, $(28.4 \pm 1.8) \%$.

Also caught A pre-existing fragility in `altitude_noise.collect` where a single unreadable file (a transient external-disk hiccup) aborted the whole PSD sample; it now skips and logs the file instead.

2026-07 — SSD reorganised; no data kept locally; figure fixes.

Policy going forward Raw data, catalogs and every analysis byproduct live only on the external SSD, never in the local repo.

Reorganisation Each discipline’s `data_root` is now laid out by maturity — `raw/{igc,raw_xml}` (untouched acquisition output), `catalog/{catalog.csv,seasons_index.csv}` (tables derived from it), `derived/` (further byproducts, today just `track_scan.parquet`; future cleaned/filtered data and analysis results get their own `processed/analysis/` siblings), with `Config` updated accordingly (new `derived_dir` property). Verified file counts before/after the move (paraglider: 186,052 tracks, 27 XMLs; hang glider: 6,716 tracks, 25 XMLs — all matched) and that both acquisition CLIs still work against the new layout. Previously-local `track_scan.parquet` copies moved to `<data_root>/derived/`; orphaned local `catalog.csv` copies (superseded by the track-based scan months ago) were removed.

Figure 4.5 (gaps) The adopted `max_gap_factor` was raised from 5 to 10 — at 5 it retained only 77.8 %/65.0 % of para/hang flights (nowhere near the near-full retention of the duration/path cuts), at 10 it reaches 86.9 %/85.1 %. Considered 20 (90.8 %/90.6 %, little further gain) and rejected it: the cut is *relative* to each flight’s own native Δt , so for the slower hang-glider loggers (up to 10 s native) a factor of 20 would tolerate interpolating across a ~ 200 s gap — long enough to span a thermal climb. Also tightened the distribution panels’ axis range to the data’s own percentiles (the old 0–200/0–0.6 range was mostly empty).

Figure 4.1 (altitude noise) Panel (a)’s difference trace now uses a colour distinct from both channels (was reusing GNSS’s red); its window grew from 10 to 30 minutes; a Nyquist-frequency line was added to panel (b) (previously discussed in the caption but never drawn). Also found and fixed a real bug in the “representative flight” selection: picking the *longest flight by elapsed time* broke on a 30-hour, 1,100-fix outlier (one corrupted timestamp gap) — fixed to select by *fix count* instead, which has no such failure mode.

Local data folder clarified The two versioned `seasons_index.csv` are the offline reproducibility anchor for the thesis’s dataset counts (the document builds with no SSD mounted), not stray local data. Their sync from the canonical SSD copy is now explicit and automatic (`refresh_seasons_index.py`, run best-effort by the pre-commit hook), so the snapshot cannot silently drift.

2026-07 — Fix-level cleaning made data-driven; parser roll-over bug; figure and data-folder cleanups.

Fix-level cleaning audited Table 4.1 is now audited on the data, like the flight-level cuts: a new figure in Sec. 4.1.2 shows the per-*fix* distributions of horizontal speed, barometric vertical speed and barometric altitude (`make_fixlevel_diagnostics_figure`, from `fix_level_distributions` over a seeded sample — millions of fixes). What justifies a fix-level bound is the fraction of *fixes* it removes, not the number of flights it touches; the per-flight *maxima* already in the scan (`max_vxy_mps` and friends) answer the wrong question,

so they do not place the cuts.

Gap handling unified Removed the absolute 60 s inter-fix time-gap from the fix-level table: it duplicated the sampling-regularity cut of Sec. 4.1.6, which is more principled (relative to each flight’s own native Δt). Neither was ever an enforced filter, so this is a spec simplification, not a behaviour change.

Parser bug fixed Found the root cause of the “30-hour, 1,100-fix” pathology worked around previously: the IGC parser’s midnight roll-over unwrap treated *any* backward step in time-of-day as a new day and added 86,400 s, so a few-second GPS glitch became a spurious multi-hour gap. Now only a drop of (nearly) a whole day counts as a roll-over, with residual backward jitter clamped to keep the series monotonic. Added regression tests for both the true roll-over and the glitch.

Figure and config cleanups The altitude-noise figure (Sec. 4.1.1) is now four panels — (a) raw channels, (b) their difference, (c) PSD, (d) fallback rate — fixing the old a/b/d labelling. The gap-diagnostics figure (Sec. 4.1.6) no longer clips its tail into an artificial edge spike, with an x-range fitted to the data’s own percentiles. Result-specific numbers (retention percentages, one-off anecdotes) were pruned from YAML/code comments — such numbers belong in the thesis, not in reusable config.

Local data folder clarified Rewrote `data/README.md` and fixed several stale references to a flat `data/seasons_index.csv` path that never existed.

2026-07 — Coordinate/geodesy figures. Redrew the coordinate-transformation figures: Figure 4.3 (WGS84 ellipsoid; geodetic vs geocentric latitude) and Figure 4.4 (the ECEF→ENU rotation, in our notation), and added the ESA Navipedia citation for the ECEF conversion.

2026-07 — Altitude-noise figure (Fig. 4.1): panels (a)/(b) redesigned. A closer look at the rendered figure (not caught by any test, since nothing there errors) found two compounding problems in the one-flight time-series panels, both traced to a single shared 30-minute window: the several-hundred-metre climb dominated the vertical scale, hiding the GNSS jitter the panel exists to show; and panel (b), subtracting only the scalar *mean* of the GNSS–barometric difference, ended up on a ± 20 m scale that looked inconsistent with the ~ 130 m gap plainly visible in panel (a).

First attempt (reverted) Shortened the shared window to 2 minutes and switched panel (b) to a *linear* detrend (matching `_welch`’s). This resolved both symptoms but destroyed something worth keeping: the GNSS–barometric difference actually drifts by tens of metres over a full flight (plausibly the barometric channel’s departure from the ICAO standard atmosphere as altitude changes) — a real, physically meaningful feature that a 2-minute window is too short to show and a linear detrend actively removes.

Final design Decoupled the two panels’ windows and mechanisms instead of forcing one compromise on both. Panel (a) keeps a short (2-minute) window but now plots barometric and GNSS on *two independent y-axes* (same numeric span, different baseline) rather than one shared axis, which previously had to stretch from one channel’s minimum to the other’s maximum — wasting on blank space exactly the room needed to compare the two channels’ roughness; with matched spans the GNSS trace is visibly more jagged than the barometric one. Panel (b) reverts to the original 30-minute window with only the *mean* removed (no detrend), so the slow drift is visible again alongside the residual jitter, with the removed mean offset (139 m for this flight) annotated directly on the panel. Both panels’ wording was generalised from “climb” to “vertical motion”, since a short window can just as well land on a glide.

2026-07 — dev/analysis moved from pip extras to uv default dependency groups.

Chasing the figure fix above, a bare `uv sync --reinstall-package soaring` (no `--extra` flags) silently reset the environment to the core-only dependency set, uninstalling `pytest/ruff/mypy/matplotlib/scipy/pyarrow` — and, on the way, hit a pre-existing uv bug (uninstall failing on packages whose RECORD file was missing) that required rebuilding `.venv` from scratch.

Root cause `dev` and `analysis` were plain PEP 621 optional-dependencies, so *any* `uv sync/uv run` invocation without matching `--extra` flags resyncs the venv down to exactly what was requested, silently dropping whichever extras the previous invocation had added — alternating flagged and unflagged commands across a session means constant install/uninstall churn.

Fix Moved both to `[dependency-groups]` (PEP 735) with `[tool.uv] default-groups = ["dev", "analysis"]`, so every `uv sync/uv run` includes them unconditionally. Neither group was ever meant to be pip-installed by an external consumer in practice (this repo is not published), so the loss of `pip install .[dev]`-style installability is immaterial; `docs` remains a real extra since it is genuinely on-demand. The docs-deploy CI workflow now opts out of the new default groups (`--no-default-groups`) to stay minimal.

Docs updated `README.md`, `docs/guide/installation.md` (dependency-groups table, pip fall-back instructions), the two reporting scripts' now-unnecessary `uv run --with matplotlib --with scipy` incantations, and the `soaring.analysis` package docstring.

2026-07 — Logbook readability; Next-steps drift; docs deploy on demand. Three unrelated loose ends, tidied together.

Logbook typography The narrative notes below were dense, single-paragraph entries — hard to scan. Reformatted the longer ones into short lead sentences plus labelled sub-points (*what changed*, *why*, *gotchas*), and dropped a stale forward-reference in the altitude-channel entry (“to confirm on a larger sample once the SSD is remounted”, superseded two entries later by an exact 186,025-flight census).

Next-steps drift The Next steps list below had fallen out of step with the thesis chapter it is meant to mirror: *Preliminary characterization* (thesis Sec. 4.2) and *Tooling* (Sec. 4.9) were missing entirely, and the whole-trajectory, phase-resolved and transport-analysis bullets were missing content the thesis already has (the propagator/return-probability/first-passage observables; the step-waiting coupling and phase-sequence memory; the stratification and bias-control plan). Re-synced against the current chapter.

Docs deploy The GitHub Pages workflow redeployed on every push to `main`, including pushes with nothing to do with the site (e.g. acquisition/analysis code with no docstring change) — not wanted. One such deploy had also failed with a generic, transient “try again later” from GitHub’s Pages API, but that was an isolated glitch, not a config problem (confirmed via the Actions log: only that one run failed out of the last ten, and only at the deploy step, after a successful build). Fix: the `push` trigger is now scoped with `paths` to what actually feeds the built site (`docs/**`, `mkdocs.yml`, `src/soaring/**`, whose docstrings `mkdocstrings` renders into the API reference), so the site still updates itself automatically but only for changes that matter to it; `workflow_dispatch` stays available for an on-demand redeploy.

2026-07 — License added; README badges; a real test-running CI. `pyproject.toml` had declared `license = MIT` since the project’s start, but no `LICENSE` file existed, and the repo had no workflow that actually ran the test suite (only the docs-deploy one).

License Added `LICENSE` (MIT): permissive, minimal text, the de facto standard for research Python code, and consistent with what `pyproject.toml` already declared — no reason to pick something more restrictive for code meant to be freely reused.

New CI: tests.yml Runs `uv run pytest` with coverage on every push to `main` and every PR. Needed no `-extra` flags to get `pytest-cov/matplotlib/scipy/pyarrow` – a direct payoff of making `dev/analysis` `uv` default groups earlier. Coverage XML uploads to Codecov as a best-effort step (`fail_ci_if_error: false`): a Codecov hiccup, or the repo not being linked there yet, must never redden the Tests badge.

README badges Added five badges (python version, Tests workflow status, Codecov coverage, docs-online link, license) at the top of `README.md`. The Codecov one needs a one-time manual step outside this repo – linking `matteodisante/soaring-anomalous-transport` on `codecov.io` (GitHub login) and, if it asks for one, adding `CODECOV_TOKEN` as a repo secret – which cannot be done from here.

2026-07 — Fig. 4.1/4.2 and Appendix B: closer reading, several fixes. A line-by-line pass over the altitude-noise figure, the fix-level figure, and the PSD appendix, prompted by direct questions about what they actually show.

Fig. 4.1 panel (a), third look The dual-axis, short-window version (see the entry above) did not read, on the page, as clearly showing GNSS is the noisier channel, so it was dropped for a shared axis over the full 30 min window instead – which showed the offset drift directly, but on request the panel went back to a dual axis once more, this time over that same 30 min window rather than a short one. Net result: each channel is independently auto-scaled (barometric left, GNSS right), so the two curves track almost on top of each other (both are the same ~ 500 m climb, just at a different baseline) – good for reading each channel’s own shape, but it means the raw vertical gap is no longer a literal reading of the offset. That number was never really this panel’s to carry: it lives in panel (b), and the noise-floor claim in panel (c), regardless of which of these three versions panel (a) went through.

Appendix B (PSD) bug Two cross-references pointed at “Figure 4.1(b)” for the PSD, which has been panel (c) since the four-panel redesign; fixed. Also added the actual flight yield of the Δt -matching step, which the text described only qualitatively: of 400 sampled flights per discipline, 212 paragliders but only 114 hang gliders pass it, because hang gliders split across several native rates while paragliders are overwhelmingly 1-second loggers – so the hang-glider PSD curve is honestly an average over that 1-second minority, not the full hang-glider population, now stated as such.

Fig. 4.2 (fix-level), added The exact sample (800 flights per discipline, 6,633,923/4,541,511 paraglider/hang-glider fix pairs) was missing from the caption; added. Also explained two things a reader flagged as confusing on sight: panel (a) (horizontal speed) shows a knee well below the adopted 40 m/s bound ($\sim 23/35$ m/s for para/hang) – read as the transition from ordinary gliding speeds to rarer but real high-energy manoeuvres (wind-assisted glides, steep spirals) rather than GPS error, hence the bound is deliberately kept above it, not moved down to it, so as not to discard genuine dynamics; and panel (b) (vertical speed)’s comb-like histogram is a quantisation artefact (integer-metre barometric altitude differenced over a 1-second native step can only return an integer m/s), not a plotting bug.

Table 4.1, added A one-sentence justification for the -100 m lower altitude bound, which read as arbitrary: it is not a claim about true terrain elevation, but headroom for the barometric channel’s own non-QNH-corrected offset (Figure 4.1 measures this directly, ~ 100 m for the flight shown there), so a genuine sea-level take-off can log slightly negative without being cut.

2026-07 — Fix-level speed bounds made per-discipline, on a much larger sample. Chasing the loosened-threshold experiment of the previous entry to a conclusion.

Sample size `FIXLEVEL_SAMPLE_PER_DISCIPLINE` raised from 800 to 15,000: 15,000 paragliders (126,723,499 fix pairs, ~ 73 s) and all 6,716 hang gliders (a full census at this population size, 37,915,334 pairs, ~ 25 s) – under two minutes total, and it resolves the sparse tail well enough

to distinguish real dynamics from a rare artefact, which the 800-flight sample could not.

The hang-glider bump, confirmed and explained At this sample size the 45–90 m/s feature is unambiguous: not a smoothly decaying tail but a distinct, much more slowly-decaying plateau (confirmed with a fine, 2.5 m/s-binned histogram – the decay rate visibly changes between 37.5 and 42.5 m/s, from a steep real decline to a near-flat plateau lasting to ~ 90 m/s). A sustained *ground* speed of 160–324 km/h has no credible airborne explanation for a hang glider: read as a logger/GPS artefact population, not rare-but-genuine manoeuvres.

Decision: per-discipline bounds, not one shared value Paragliders and hang gliders have different performance envelopes, confirmed by their different tail shapes – a shared bound is either too loose for one or clips real dynamics of the other. `FixLevelThresholds.max_horizontal_speed_mps`, `max_vertical_speed_mps` are now `dict[str, float]` keyed by discipline (`configs/preprocessing.yaml`, a nested mapping; a "sailplanes" entry can be added later with no code change). Vertical speed: 12 m/s (para) / 14 m/s (hang).

Horizontal speed bound, revised again: 40 was still “slowing”, not “stopped” First adopted at 25/40 m/s (para/hang), where the fine-binned decay rate first visibly changes. On review, with the panel’s x -range widened (next entry) to actually show the far tail, that point turned out to still be on the shoulder of the decline, not the flat artefact itself: the bin-to-bin ratio only settles at ≈ 1 (a genuine plateau, decay actually stopped) around 55 m/s for hang gliders, and the analogous point for paragliders is ≈ 45 , not 25. The 25/40 m/s pair was therefore cutting into the tail-end of real (if rare) dynamics – a wind-assisted glide, a steep spiral – rather than sitting cleanly past it. Revised to **45 m/s (para) / 55 m/s (hang)**, both now Table 4.1’s and the config’s adopted values.

Figure and table updated `make_fixlevel_diagnostics_figure` now draws one dashed cut per discipline, colour-matched to its histogram, with a per-discipline removed-fraction annotation (previously one shared cut and fraction). Table 4.1 gained a row per discipline for the two speed bounds.

–100 m altitude bound, empirical check (unchanged, kept for the record) Computed the per-flight mean GNSS–barometric offset over 1500 flights per discipline (broader than the single Figure 4.1 example): median ≈ 90 m, 25th–75th percentile roughly 43–135 m. This is a *whole-flight* average, and the offset itself grows with altitude (Sec. 4.1.1), so it likely *overstates* the offset specifically near the ground; confirms the order of magnitude of –100 m without pinning it down precisely. No error found in Figure 4.2 itself.

x -range widened for (a)/(b) The panels’ range was cropping the view almost exactly at the cut (the 99.9th percentile of the pooled sample sits right there, by construction), hiding how much – and what – the bound leaves out. Panels (a)/(b) now extend to each discipline’s own 99.99th percentile (computed per discipline, not pooled, so hang gliders’ heavier tail is not swamped by paragliders’ larger sample); panel (c) (altitude) keeps the original, narrower range on purpose – its cuts sit far out in an otherwise tight, unimodal distribution, and widening it the same way mostly shrinks the informative part of the panel.

Panel (b)’s non-empty valleys, explained and checked The comb (spikes at integer m/s) was expected; that the valleys *between* integers were not exactly zero was not, and looked like a bug. It is not: not every flight samples at 1 s (Figure 4.7 – hang gliders spread across several native rates), and differencing over a longer native step returns a finer-grained set of values (e.g. multiples of 0.2 m/s at a 5 s cadence) that fill the gaps once every cadence is pooled. Checked directly, split by native rate: among 1 s-native fixes, $\sim 99\%$ land within 0.05 m/s of an integer (the comb, essentially pure); among every other native cadence, only 34–42% do – the other 58–66% is exactly what shows up between the teeth.

Panel (b)’s jagged comb, fixed Separately from explaining *why* the valleys were not empty: the comb itself looked bad – an arbitrary, fine (linspace) grid of bins straddles the integer spikes inconsistently depending on where each bin edge happens to fall, which reads as noisy

up-down jitter rather than the smooth decaying envelope the data actually has. Fix: one bin per integer, centred on it (edges at half-integers) rather than an arbitrary edge grid – `make_fixlevel_diagnostics_figure` now takes an `integer_aligned` flag, set for panel (b) only (`vxy` is not similarly quantised). The panel is now a clean, monotonically-decaying curve; the removed-fraction numbers are unaffected (computed from the raw data, not the histogram).

Vertical speed reverted to one shared bound: 13 m/s Unlike horizontal speed, paragliders’ and hang gliders’ $|v_z|$ distributions sit close enough together (panel (b)) that a per-discipline split buys nothing; a single 13 m/s bound (between the previous 12/14 split) covers both, with one shared grey dashed line and one pooled removed-fraction annotation, mirroring how the altitude band already works. `FixLevelThresholds.max_vertical_speed_mps` reverted from `dict[str, float]` to a plain `float`; the figure’s two shared-cut helpers (band and single-cut) were merged into one `_shared_cut_panel` that handles either, so panel (b) still gets the `integer_aligned` treatment above.

A LaTeX limit hit, and fixed for good: very long captions Rebuilding after the above changes hit ! `Dimension too large` / “I can’t work with sizes bigger than about 19 feet”, deep in `\@makecaption`. Bisected the Figure 4.2 caption by half repeatedly to isolate it – not a typo (braces and math-mode $\$$ were balanced throughout), but the caption’s sheer length (~3900 characters): plain L^AT_EX’s single-line-fit check for captions measures the *whole, unwrapped* caption as one hbox before deciding whether to wrap it, and past a few thousand characters that measurement alone overflows T_EX’s ~19-foot maximum dimension, even though the wrapped result would render fine. Loading the `caption` package does the same check by default (to centre genuinely short captions), so `singlelinecheck=false` was needed on top of it. This is a general limit of this document now that several figures carry long, data-justified captions – not specific to this one – so the fix (`main.tex` preamble) covers all of them going forward, not just Figure 4.2.

2026-07 — Figure 4.7 (native sampling interval): discrete data, discrete plot. Raised independently: the density in Figure 4.7 rose above 1, and native Δt is a discrete quantity, so why plot a density at all. Checked on the full census: Δt lands on an exact whole second for 99.9% of paragliders and 100% of hang gliders – essentially a categorical variable (one of a handful of configured logger rates), not a continuous one. The previous plot used arbitrary 0.5 s-wide bins on this discrete data (`np.arange(0.25, 8.75, 0.5)`): most paragliders’ mass concentrated in the single bin straddling $\Delta t = 1$, and `density=True` divides a large fraction by that narrow bin width, which is exactly how it exceeded 1 – a real number, not a bug, but a genuinely confusing one for data this concentrated. Fixed the same way as the vertical-speed comb: bins one full second wide, centred on each integer (1, ..., 10, plus a ≥ 11 bin for the long thin tail), so each bar reads directly as “fraction of flights at exactly this rate.” With unit-width bins `density=True` is numerically identical to that fraction (dividing by a width of 1 changes nothing), so the axis is now labelled *fraction of flights*, not *density* – accurate either way, but only one of the two labels says what is actually being counted. Real distribution, now stated precisely in the caption instead of the vaguer “1, 2, 5, ≥ 8 ” before: paragliders overwhelmingly (69%) at 1 s; hang gliders spread across 1 s (30%), 5 s (27%), 10 s (15%), 2 s (13%), and a long thin tail beyond 11 s – a fuller picture than the old caption’s “up to 8 s,” corrected in Sec. 4.1.6 too.

2026-07 — Figure 4.6 (gap diagnostics), panel (a): dropped the log axis. Same family of question as the two entries above, this time about the largest-gap-to- native- Δt ratio: is it also discrete, why a density, and (specifically here) is a log x -axis actually needed.

Not as clean-cut as native Δt Checked directly: only 86% of paraglider and 63% of hang-glider values are exact integers (thousands of distinct values in all, tail out to $> 50,000$) – a missed fix or two at the native rate is the single most common cause of a large gap (hence

the pull towards small integers), but plenty of gaps are not a clean multiple of it. So, unlike Figure 4.7, this is not purely categorical; integer-only bins would have thrown away a genuine continuous component.

The log axis, though, was not earning its keep Panel (a) only ever displays the *bulk* of the distribution (up to `gap_hi`, a modest $\sim 1\text{--}20$ here) – the heavy tail is deliberately left to the retention panels (c)/(d), which already sweep the full range on a log axis, justified there since the cut spans orders of magnitude. Over the narrow, already-cropped range panel (a) actually shows, a log axis buys nothing and actively hides structure: log-spaced bins are extremely narrow near 1, so the large mass sitting at small integers – a real, common failure mode, not noise – gets compressed into a sliver.

Fix Panel (a) switched to a linear axis with linear, evenly-spaced bins over the same `gap_hi`-bounded range. This did not just fix the axis choice: it revealed a real pattern that the log/narrow-bin combination had been washing out – visible peaks at small-integer ratios, most clearly for hang gliders (around $k = 3, 6, 10, 15$), consistent with the “missed native-rate fixes” mechanism. Panel (c)’s log axis is unaffected and still the right choice there, for the opposite reason: it has to show a cut ranging over orders of magnitude.

2026-07 — Why gap ratios aren’t integers: stratified by native rate. Follow-up on the panel (a) entry above: if a missed native-rate fix is the single most common cause of a large gap, why is the largest-gap ratio an exact integer only 86%/63% of the time, not effectively always? Stratifying the full census by native Δt answers it.

$\Delta t = 1\text{ s}$ is trivially always integer Every IGC timestamp already has whole-second resolution, so the largest gap itself is always a whole number of seconds; dividing it by $\Delta t = 1$ cannot produce anything but an integer. Checked directly: 100% for both disciplines at this rate. Since 1 s is also the dominant paraglider rate (69% of flights, Figure 4.7), it alone pulls the paraglider-wide fraction up to 86%.

Slower configured rates are only about half integer At $\Delta t = 2/5/10\text{ s}$ the integer-ratio fraction is 60%/58%/51% for paragliders and 34%/61%/39% for hang gliders – no rate stands out, none close to 100%. The mechanism from the panel (a) entry still holds (a missed fix is still the main driver of a large gap), but the gap itself is a real elapsed duration – GPS reacquisition time, logger buffering, a signal dropout – set by whatever caused the interruption, not by the logger’s configured cadence. Nothing requires that duration to land on a multiple of Δt ; it does so about half the time by chance, not by construction.

So the headline 86%/63% is a mixture effect Paragliders are predominantly 1 s loggers (trivially integer), hang gliders are spread across several slower rates (mostly not) – the two overall percentages mainly reflect which rate dominates each fleet, not some fleet-wide regularity in how gaps occur.

2026-07 — Figure 4.6, panel (b): tightened the missing-fraction axis. The x -axis ran out to roughly twice the adopted cut (0.10), leaving the right half of the panel empty: the 90th percentile of the pooled missing fraction sits at ≈ 0.09995 , essentially at the cut itself, so nothing of substance lives beyond it. Changed the margin past the cut from $2\times$ to $1.1\times$, i.e. `miss_hi = max(p90, cut * 1.1)`, giving a 0–0.11 range that still clears the 90th percentile with a little headroom but no longer wastes half the panel on empty axis.

2026-07 — Figures now stay where they are placed in the source. Plain L^AT_EX’s float algorithm was free to defer any figure past its ideal spot to the next page it judged to have room, and with several large figures placed close together in Sec. 4.1 it was pushing some many pages away from the text discussing them – annoying for anyone reading top to bottom rather than jumping via cross-references. Added the `float` package and switched every `figure/table`

placement specifier (previously `[t]/[!]`) to `[H]` (“here, literally”), which pins each one to its position in the source, full stop. No layout regressions on a full rebuild: text and figure land on the same page throughout, including the two cases checked directly (Figures 4.6 and 4.7, immediately after the paragraphs that introduce them).

2026-07 — Thesis restructured: narrative body, implementation appendix. The body had grown heavy: captions and prose mixed physical/statistical reasoning with code module names, exact config values, sample sizes and computation times (worst case, Fig. 4.2’s caption, 3898 characters). Restructured so the body reads like an ordinary thesis and every implementation-level fact moves to one place, cross-referenced tersely from where it used to sit.

New back-matter part. “Implementation and Computational Details,” a document-within-the-document at the very end: unnumbered part divider (a numbered “Part I” would have implied a nonexistent “Part II”), its own local “Contents of this part” listing (`titletoc`), and its own independent chapter numbering (plain arabic, reset – continuing the A/B appendix lettering would have made it read as “just more appendices”). Mirrors the body chapter by chapter: Overview, Data acquisition, The dataset (incl. pre-processing and preliminary characterization), Next steps.

One macro, used throughout. `\impldetails{impl:x}` renders “(implementation details: Sec. X)” wherever a fact was trimmed out of the body; `impl:*` labels mirror `sec:*` 1:1.

Chapter 3/4 boundary moved. Trajectory pre-processing and preliminary characterization are dataset description, not planned analysis – moved from “Next steps” into “The dataset” (which now runs to Sec. 3.6); “Next steps” starts at “Observables on the whole trajectory.” Five cross-references that would otherwise have become circular self-references were updated in the same pass.

Every pre-processing figure/table slimmed. Fig. 4.2’s caption: 3898 to ~1500 characters, keeping the artefact argument (hang-glider plateau, physically-absurd km/h translation) and moving exact sample sizes/timings/bin mechanics. Table 4.1, Figs. 4.1/4.6/4.7 similarly. `B-psd.tex`: three clauses extracted from “Welch’s method” (segment length, `scipy` name, overly specific detrending framing), its whole “Computational details” section relocated, one genuine external-validity caveat (hang-glider PSD averages a 1 s-logger minority) preserved in “Reading the altitude spectra.” `02-data-acquisition.tex`: CLI/subcommand names, atomic-write/retry mechanics, env-var names relocated; two hardcoded approximate flight counts fixed to use the `\Stat*` macros instead (drift risk).

One Python edit. The two season-table captions (`scripts/reporting/generate_stats.py`) had their provenance clause replaced with the macro; regenerated, `ruff` clean.

Verified throughout: every `impl:*` label matches an `impldetails` call (no dangling cross-reference), full rebuilds clean at each stage (no undefined references), official build script (`scripts/build_docs.sh thesis`) green.

2026-07 — Fix-level cleaning redesigned: three detectors, unambiguous removal.

A critical re-reading of the fix-level cleaning, asking whether the adopted design (Table 3.6.2 / old 4.1) is actually reliable, exposed real gaps and turned the flat list of absolute bounds into a design organised by *how* each defect presents. Cleaning only, not the flight-level filtering.

The gap that mattered “Reject the fix” was ambiguous on the one question that decides everything: when the speed between two fixes is too high, *which* of the two is removed? A speed lives on the segment between fixes and cannot say. And a fixed global speed bound is only a teleport filter — it passes the moderate spikes (a 30 m/s step for a 12 m/s paraglider) that dominate the differentiated velocity/acceleration, and is blind to a frozen GPS lock (repeated position, zero speed, clock still running), which mimics a long CTRW waiting time and would plant false trapping events in the very tail the anomalous exponent is read from.

Guiding principle Cleaning must not bias the transport statistic it precedes: assume only that the trajectory is smooth at the sampling scale, never a model of the motion. This rules out the tempting “best” upgrade, a Kalman / robust motion-model smoother — its constant-velocity or diffusive prior would write the studied behaviour into the data.

Three detectors, by context *Off the trend* (spikes): a robust local-outlier test (Hampel — median centre, MAD scale, threshold in robust σ , plus an absolute floor $\varepsilon_{\min}$), a per-*node* quantity, which names the offending fix directly and dissolves the which-fix ambiguity. *On the trend* (structural): duplicate time $\rightarrow$ centroid; non-roll-over backward time $\rightarrow$ delete; frozen-lock run $\rightarrow$ mark as gap. *Out of range*: the old absolute bounds, kept as the interpretable backstop the figure already audits.

Two removal semantics Split apart what “remove” meant: a corrupt position/time deletes the node (gap bridged later at resampling); a corrupt altitude invalidates only the altitude channel, keeping the horizontal position — so the primary (horizontal) MSD loses nothing to a barometric spike. The old table coupled the two, throwing away good (x, y, t) for an altitude fault.

Frozen lock $\rightarrow$ split, not bridge A lost-lock run is marked as a data gap and the flight is *split* there into independently analysed segments — not interpolated across (which invents motion) nor dropped whole (which wastes the good sides). Folded into the uniform-sampling step, which already floated segmentation as a refinement.

Validation, both directions The existing “removes a negligible fraction of fixes” argument only bounds false *positives*. Added two checks for the other direction: an injected-defect / recovery test (planted spikes/freezes/time-glitches of known size give the detection rate, i.e. the false negatives, and calibrate the parameters) and a downstream-invariance test (MSD exponent and waiting-time tail versus cleaning strength; adopt the plateau where they stop moving — the evidence that error, not signal, was removed).

Where it lives Narrative in Sec. 3.6.2 (Table regrouped by family with an explicit action column, fix-level figure reframed as the backstop audit); full working definitions — Hampel parameters, frozen-lock thresholds, single-pass $O(N)$ anchor scan, integrity gate, validation protocol — in the dataset chapter of the implementation appendix. The cleaning routine and its new config keys $(w, k, \varepsilon_{\min}, \varepsilon, \tau_{\text{freeze}}, f)$ remain to be built; the blueprint (`preprocessing-pipeline.md`) and its open items were updated to match. Thesis compiles clean (50 pp, no undefined refs, no new overfull boxes).

2026-07 — Frozen-lock witness ranked; the cleaning spec quantified. A consistency audit of the frozen-lock rule found a real flaw: the V/zero-alt flags sat in a plain or-branch and could overrule a climbing barometer, while the IGC standard says V certifies the GNSS *altitude*, not the horizontal fix.

Fix The witness of `eq:frozenlock` is now ranked and case-split by `alt_source`: on a barometric flight, baro-flat or byte-identical repeats (the only baro-independent override, since no moving receiver rewrites identical bits); V/zero-alt decide only on GNSS-fallback flights. All four contract locations updated together (thesis body, implementation appendix, blueprint, `preprocessing.yaml` comment).

Guarantee quantified A real climb misses the cut on two independent instruments: $2R \approx 40\text{--}100\text{ m}$ against $\varepsilon \approx 10\text{--}15\text{ m}$, and at least 60 m of barometric gain over τ_{freeze} against a few metres of δ_z . The one assumption (a working barometric channel) is now stated, with the residual corner cases confined to the GNSS-fallback minority.

Presentation Hampel identifier named and cited in the body; the speed bound reframed from “backstop” to corroboration gate, matching the appendix; `tab:cleaning` regrouped with an invariant row; the seven detector parameters collected in a working-defaults table (`tab:fixlevel-params`); thresholds stated inline where each decision is made (the prose said

“a few Δt ” where the adopted bound is $10 \Delta t$).

2026-07 — Acquisition folded into the dataset chapter. Data acquisition is no longer a chapter: it opens “The dataset” as Sec. 2.1 (the file is `\input` there), its implementation mirror folded into the dataset mirror the same way, and the introduction roadmap rewritten; the thesis now has three chapters. In the same pass the pre-processing order (i)–(vii) became an enumerate mapping the chapter, the trimming rule became an equation ($t_{\text{on}}/t_{\text{off}}$ as min/max of sustained $v_{xy} > v_0$ windows), and the implementation appendix’s worst walls of text were rewritten in recipe form (Hampel test as numbered steps, detector order and the five validation checks as lists).

2026-07 — Voice pass: the document must read human. The prose had accumulated a generated-sounding register, partly from the AI-assisted sessions themselves: narrativized headings (“One rule: clean without biasing the transport”), aphorism setups, meta-commentary (“the ranking is the point”), a courtroom metaphor sustained across the cleaning section, and dense em-dash chains. Swept the whole document: headings are now nominal (Design principle, Removal policy, Frozen-position runs), the cleaning narrative rewritten plainly with content and numbers unchanged, the metaphor kept only where it is notation (the witness W), a duplicated variometer sentence removed, prose dash pairs roughly halved. Appendix A already read plain and was left untouched.

2026-07 — Split semantics pinned down. Asked directly whether splitting a flight at a long gap can bias the analysis. The split itself is the right call (interpolating fabricates motion, dropping the flight wastes data), and ensemble observables are safe by construction: a segment keeps its parent’s origin and clock, and GPS positions are absolute, so $\mathbf{r}(t)$ after the gap is valid and the flight merely contributes nothing at the missing lags. Three cautions were previously unstated and are now in the spec: phase durations truncated at a segment boundary are censored, not complete (else every split reads as a spuriously short wait and thins the $\psi(\tau)$ tail); a segment is an aged observation window, not a fresh flight, for the time-averaged MSD; and the flight-level cuts are *not* re-applied to segments — the earlier draft said they were, which would have preferentially discarded the data behind late-flight gaps — replaced by a minimal segment-duration gate (config key with the routine).

2026-07 — Second split audit: the gap bound gains an absolute scale. Two questions from a design review, before implementation. (1) The gap bound was purely relative ($10 \Delta t$): a 3 s logger could bridge 27 s, a 10 s logger 100 s, while the harm of a bridged gap is set by the motion’s timescales, not the logger’s (an interpolated stretch is a straight line; near the thermalling period it can replace a full circle and read downstream as a glide). The split bound is now $\min(10 \Delta t, 20 \text{ s})$: the cap is at least twice the slowest common cadence (one missed fix never splits), at most about two thirds of a thermalling period, and well below τ_{freeze} (an excised frozen run is split, so a native gap of that length cannot be bridged). Added as `max_gap_seconds` to the YAML and the `SamplingThresholds` dataclass, with the config test extended. (2) How a split works mechanically: the ENU conversion precedes it, so both segments already carry parent-referred coordinates and time; a split assigns a `segment_id` in the `fixes` table, per-segment records go to a `segments` table — no separate files, and the ensemble MSD reads $\mathbf{r}(t)$ wherever a segment covers t . Also flagged: trimming strips only the *outer* ground phases, so a mid-flight landing and relaunch in one log would read as an implausibly long wait; an interior-ground guard (sustained near-zero v_{xy} of order ten minutes, treated like a long gap) is planned with the implementation.

2026-07 — Third audit: interpolation budget, interior landings, per-segment gates.

(1) Interpolation error over a bridged gap is governed by the unresolved manoeuvring, not by the interpolant order: during a glide the error over 10–20 s stays at metres (below GPS noise), while during circling a gap of g seconds spans an arc of $2\pi g/T$ that no interpolant can recover — at $g = 10$ s the chord-arc deviation is already $\approx 0.5R$ (10–25 m), at $g = 20$ s up to $2R$. This confirms the two-scale bound and adds a refinement: the cap is floored at $2\Delta t$ ($g_{\max} = \min(10\Delta t, \max(20 \text{ s}, 2\Delta t))$), so the thin tail of very slow loggers also tolerates a single missed fix (a bridged gap there inserts at most two grid points). Every filled grid point is now flagged **interpolated** in **fixes**, so phase statistics can treat reconstructed stretches with the same caution as segment boundaries. (2) The interior-ground detector is specified, built like the frozen-lock rule (cut only on corroborated evidence, since a wing soaring into wind holds near-zero ground speed while flying): $v_{xy} < v_0$ sustained for T_{ground} (order ten minutes) *and* barometric altitude flat within a few metres; a detected stint is excised and handed to the split machinery. (3) Two gate corrections: the integrity gate counts over the airborne window only (ground junk that trimming removes anyway must not condemn a flight), and the missing-fraction gate acts per segment (a bad segment is dropped alone). Savitzky–Golay edge handling pinned (`mode='interp'`, per segment).

2026-07 — Fourth audit: the gap cap measured on the census; blueprint realigned.

With the SSD mounted, the two-scale gap bound was measured on the cached census scans (paired Δt /largest-gap per flight, thresholds read from the config, never hard-coded): the combined bound leaves 15.3% of paraglider and 40.5% of hang-glider flights with at least one split, against 11.8%/14.5% under the relative rule alone. The increase sits at the slow cadences ($\Delta t \geq 2$ s), where a 20 s hole is only a few missed fixes. A split is cheap for ensemble observables but fragments within-segment statistics, so the cap stays a working value and enters the downstream-invariance sweep; the shares are generated census macros (`StatScan*GapSplit*`) quoted by the thesis. Blueprint realigned with the thesis: the stale “no second absolute bound is needed” sentence removed, `segment_id/interpolated` added to the **fixes** schema, a **segments** table added, Savitzky–Golay per-segment `mode='interp'` recorded, open items refreshed (δ_z , `min_segment_duration`, `T_ground`, suspect-wait flag), and the header now states the working rule: implement from the blueprint, thesis normative on conflict. Sub- T_{ground} stints are now *flagged* rather than merely auditable: the suspect-wait flag (near-zero v_{xy} with flat barometric trace, any duration) feeds $\psi(\tau)$ sensitivity re-fits. README refreshed: `soaring.analysis` exists, and the logbook is committed with the repo, not git-ignored as the old text claimed.

2026-07 — Fifth audit: implementer’s walkthrough; C2 de-walled. Traced one IGC file through every step of the spec as an implementer would, asking at each point “what exactly does the code do here?”. Four real logical bugs surfaced, all fixed across the contract:

Backward-time attribution “Delete the fix with the corrupt time” did not say which one; a running-max rule would let a single forward-jumped clock condemn every later fix. Now: minimal removal (complement of the longest increasing subsequence).

Position discontinuity A re-acquisition offset (track jumps and stays) fits neither the spike model (no bounded block restores the trend — deletion would cascade) nor the frozen-lock model. Now its own row in `tab:cleaning`: split at the step, like a long gap; both sides are genuine, only the transition is unknown.

interpolated was undefined under resampling Every grid point is technically interpolated once a flight is resampled. Now: measured iff a fix lies within half a native step of the grid time; the same tolerance defines `was_resampled`.

Savitzky–Golay window collapse $w = \text{odd}(\tau_c/\Delta t)$ falls below a cubic fit’s minimum for slow loggers (Δt of order τ_c or above). Now floored at $w = 5$, with the over-smoothing caveat

stated and those flights already stratified out of Δt -sensitive observables.

Smaller closures: δ_z got a number (≈ 5 m); Hampel needs a minimally populated window (order five fixes, else absolute bounds only); the frozen-lock Δz witness is median-based so an uncleaned altitude spike cannot fake a climb; suspect-ground stints are stored as per-flight intervals (phases do not exist yet at pre-processing time); a flight with no sustained-speed stretch is discarded with reason; duplicate uploads added to the blueprint's open items. Readiness verdict: the pipeline is implementable end-to-end today up to smoothing; the three Savitzky–Golay τ_c remain the one true blocker (placeholders pending the ENU PSD study; procedure and code path exist). Remaining C2 walls of text (uniform-sampling thresholds, segment handling, flight-filter diagnostic sample, PSD band reading, flag-to-action) rewritten as labelled lists, same content.

2026-07 — Sixth audit: every hyperparameter pinned; the flight filter re-measured.

Goal: implementable through step (vii) included, no ranges left.

τ_c measured, blocker closed New script (`estimate_savgol_timescales.py`): Welch on E, N (local equirectangular) and on the altitude channel each flight uses, seeded 1 Hz sample (908/719/189 flight-channels). All three knees sit at $f_c \approx 0.2$ Hz $\Rightarrow \tau_c = 5$ s: every channel's typical floor is IGC quantization (measured 0.60/0.26 m²/Hz against theoretical 0.57/0.17), not receiver noise. The baro/GNSS asymmetry lives in the spectrum's level, not its knee; even p90 spectra move it by less than a second. Placeholders replaced with measured values.

Flight filter was cutting the wrong thing Census forensics: above 40 min, path < 10 km catches 6 flights of 184,583 while the adopted 30 km cut removed 5,229 genuine localized flights (median extent 7.4 km) — the most-trapped tail, whose exclusion would bias the transport toward travellers. The declared minimal-threshold strategy itself dictated lowering: `min_path_km` 30 $\rightarrow$ 10. Also found and fixed: fifteen “flights” of 16–166 h (loggers left running, paths to 10,000 km) now cut by an upper duration bound (16 h); a few hundred retained records with total altitude range under 100 m (vehicle logs or stuck sensors) now caught by an altitude-activity floor (75 m) plus a baro-liveness check at channel choice (range < 30 m $\Rightarrow$ fall back to GNSS — a stuck baro previously passed as a baro flight with $v_z \equiv 0$).

All working ranges pinned Hampel ($w = 20$ s, $k = 5$, $\varepsilon_{\min} = 15$ m, min. window 5 fixes), frozen-lock ($\varepsilon = 15$ m, $\delta_z = 5$ m, $\tau_{\text{freeze}} = 60$ s), $f = 0.10$, segment gate 90 s, interior-ground (600 s, 5 m): all now YAML keys with dataclass fields and config-test assertions. The injected-defect calibration owns the final values.

Figures regenerated The pre-processing diagnostics were rebuilt from the SSD with the new thresholds (the path-cut line moved to 10 km).

2026-08 — Global transport measured, and three ways it could have been wrong.

The displacement statistic about take-off turned out to be unusable for an exponent, for three reasons that compound: each flight carries its own course, the CFD scores closed tasks so 56% of flights fly a triangle and come home, and the strata of Sec. 2.8 are incompatible at 40–62%. Subtracting the course is not the repair it looks like — the natural estimate uses the record's own endpoint, which turns the series into a bridge and returns $\alpha = 1.35$ on synthetic data of true 1.50, with no safe lag range.

What works is filtering rather than subtracting: an order- p difference annihilates every polynomial of degree $p - 1$, so no drift is estimated and none can be estimated badly. Over 60–2000 s the second-order filtered variation gives $\alpha = 2.02 \pm 0.10$ for paragliders and 1.94 ± 0.13 for hang gliders, the two disciplines agreeing within their uncertainties. Removing the course *raises* the exponent by 0.27, which is the opposite of what a drift does and follows from the closed tasks suppressing net displacement.

The estimator survives what the ensemble MSD does not: cutting the archive by cadence, wing class, season and task type moves the exponent by 0.02, 0.06, 0.03 and 0.10, where the

same strata departed from the pooled ensemble MSD by 40% and 62%. Chapter 3 is written on that, and reports one effective exponent over a declared window rather than a regime count — breakpoint selection returns one break and surrogates return one 16% of the time.

Four errors of my own were caught by writing the tests rather than by reading the code: a factor-two normalisation error in the synthetic fractional-noise generator, invisible to any exponent test; a false-positive rate that was circular by construction, comparing a bent curve against copies of its own curvature; a “clean range” criterion that tightened as the archive grew, so it measured the sample size; and 101 phantom test passes from five Finder-style duplicate files. The reading $\alpha = 2$ as “ballistic, nothing to see” was also wrong and was caught by calibration: a locally straight process reads 3.02 on this estimator and a straight line reads zero, so 2 means persistent and still rough.

2026-08 — The pipeline built, and the first MSD audited. Stages (i)–(vii) implemented end to end and run over both archives: 155,788 paraglider flights retained of 186,052 (83.7%) and 6132 of 6716 hang-glider ones, written as `fixes/segments/flights_meta` and checked by `verify_dataset.py` (every invariant holds) and `check_reproducible.py` (500 sampled flights re-run through the code and identical). One bug found by the verifier’s own blind spot: reading Parquet row groups instead of whole flights split every flight straddling a boundary, over-counting flights by 0.7% and truncating exactly the longest segments; `stream_flights` is the single correct reader that resulted.

Then the two MSD estimators, and an audit of what they mean. The audit is a second streaming pass keeping each flight’s position at every lag, because the questions worth asking of an averaged curve are reductions of a per-flight quantity that no longer exists once the average is taken. What it found changed how the result is reported: the ensemble curve is *not* a power law (residual 15.2% and 31.2% rms against 3.8% and 3.6% for the time-averaged one), its local slope runs 1.31–2.21 and goes super-ballistic, and that shape survives fixed cadence, fixed duration and the removal of the common heading. A slope above 2 needs a ballistic component switching on at a finite time, and the first crossing of 5 km from take-off has median 1911 s — the scale at which the overshoot peaks. So the exponent is read off the time average and the ensemble curve yields a timescale; the aging reading, which compared two fitted slopes, is withdrawn, since one of them is not an exponent.

Sec. 2.8 built with it: retained-ensemble statistics, take-off maps on real Natural Earth coastlines (three frames — metropolitan France, La Réunion, and the world for the rest), and two results that constrain Ch. 3. The ensemble is anisotropic ($\langle E^2 \rangle / \langle N^2 \rangle = 0.80$, flat over three decades), so the 1-D marginal cannot stand in for the 2-D propagator; and the strata split by wing class (40%) and terrain (62%) but not by season (12%), so the archive may be pooled over time and not over class or terrain. Verified: 383 tests, macro contract green, `regenerate.sh` end to end reproducing every quoted number, thesis compiling at 99 pages with no errors.

2026-07 — Seventh audit: global consistency sweep. Closing pass over every document at once, hunting statements left behind by the day’s changes rather than new design flaws. Thirteen stale spots found and fixed: ranges still quoted where values are now pinned (k , w , ε , τ_{freeze} , in body and appendix); “joins the config when the routine is built” where the keys now exist; the blueprint carrying two open items that contradicted each other about the same parameters; “two track quantities” where the filter now has three criteria; the repro-table caption counting two figure scripts where there are four; the logbook’s own Next-steps still calling the timescales placeholders. One reconciliation of substance: the body’s smoothing section predicted asymmetric windows (smooth barometer, noisy GNSS) and the measurement contradicted it — the knees coincide, the asymmetry lives in the spectrum’s level — so the text now states the expectation, the measurement, and why the per-channel conditioning stays. The measured-timescales script joined the reproducibility table. Verified end to end: no stale grep hits, mkdocs builds, 78 tests

pass, thesis and logbook compile with no errors, no undefined references, no overfull boxes.

2026-06 — Documentation. Project docs are built with MkDocs and published on GitHub Pages. Initially deployed with `mkdocs gh-deploy`, which accumulated one built-site commit per deploy on a `gh-pages` branch; migrated to the official GitHub Pages *artifact* workflow (build → upload artifact → deploy), removing the branch and keeping the git history clean.

Operational note. In this repo `uv sync` occasionally fails to (re)install the editable package; `uv sync --reinstall-package soaring` fixes it.

3 Next steps

These mirror the two forward-looking chapters of the state document (`thesis/sections/04-global-transport` and `thesis/sections/05-flight-phases`) and are kept in sync with them, section by section. The near-term plan is to analyse the data *in parallel* with studying the segmentation, since many observables can be built before any phase splitting.

- **Data sources.** FFVL paraglider and hang-glider CFD data are collected, cataloged and verified. Extend to further sources – notably WeGlide for sailplanes (a higher-quota API request is pending; the default rate limit makes a bulk export infeasible) and possibly XContest – merging into one catalog and stratifying by source. Refresh the per-discipline `seasons_index.csv` via `build-catalog` so the document statistics stay current.
- **Pre-processing (thesis Sec. 2.7).** The parser is in place (`soaring.analysis.igc`) and the pipeline is designed, thresholded (`configs/preprocessing.yaml`) and audited on real data: fix-level cleaning (great-circle speeds), v_{xy} -based trimming, data-driven flight-level filtering (duration + path length), the ENU conversion (barometric vertical), native- Δt uniformity, and Savitzky–Golay smoothing/differentiation. **Built and run** (2026-08): the seven stages are production code under `soaring.analysis.preproc`, driven by `scripts/preprocess.py` into the three Parquet tables, unit-tested, and checked against the invariants of Ch. 2 by `verify_dataset.py` and against the code itself by `check_reproducible.py`. The Savitzky–Golay timescales are measured (all knees at ≈ 0.2 Hz, so $\tau_c = 5$ s). Remaining: the calibration battery for Hampel, frozen-lock and the integrity gate, via injected defects and downstream invariance.
- **Preliminary characterization (Sec. 2.8). Done** (2026-08). Trajectory statistics on the retained ensemble; take-off maps on real coastlines, three frames (metropolitan France 94.7%, La Réunion, the world for the rest). Two results that constrain the analysis rather than merely describe it. Isotropy *fails*: $\langle E^2 \rangle / \langle N^2 \rangle$ sits at 0.80 for paragliders across three decades, so the 1-D marginal may not stand in for the 2-D propagator and the components are analysed separately. And the strata are *not* compatible by wing class (40% typical departure from the pooled MSD) or orographic group (62%), though they are by season (12%) — so the archive may be pooled over time, not over class or terrain. Remaining: displacement maps.
- **Global transport (Ch. 3, no segmentation). Measured** (2026-08). The chapter now rests on the filtered variation rather than on displacement from take-off: $\alpha_2 = 2.02 \pm 0.10$ and 1.94 ± 0.13 over 60–2000 s, stable under every stratification the archive allows (cadence, wing class, season, declared task; spread at most 0.06 against the 40–62% the ensemble MSD showed on the same strata). **Both MSD estimators are built and audited**: a single exponent describes the time-averaged curve ($\alpha_{TA} = 1.732$ and 1.688 , residual under 4%) and only summarises the ensemble one (15.2% and 31.2%), whose local slope crosses

over — super-ballistic at the top — where the population leaves its launch area. So the ensemble curve yields a timescale rather than an exponent, and the ergodicity test cannot take the form of a gap between two fitted slopes; it moves to the EB parameter and the phase-resolved statistic, which compare distributions. **Beyond the second moment**, also built: the moment spectrum is straight (ν spread 0.009), so *not* a Lévy walk and no strong anomalous diffusion; the non-Gaussian parameter is +0.04 for paragliders and +0.18 for hang gliders, against a Gaussian consistent with zero at large samples and 0.55 for a Lévy walk of index 1.5 — so monofractal but not Gaussian, and four times less so on the slower wing; the velocity autocorrelation has a non-integrable tail, $\gamma = 0.57$, which is what reconciles a memory already below $1/e$ at 60s with correlated increments out to 2000s; and the persistence-run tail index moves from 4.41 to 2.24 as the sinuosity threshold is scanned, so there is no distinguished leg duration — a prediction about any segmentation. **First-passage times were tried and do not serve**: $T(L)$ is measured from take-off, so it carries the launch geometry into the space domain and its exponent moves between 1.51 and 2.03 across the available radii — the ensemble-MSD crossover seen along the space axis. **The propagator of the increments is built** (2026-08) and is the only estimate of H here that comes from neither a moment nor the common origin: quantiles of $|\Delta x|$ read off histograms, one streaming pass, twelve minutes. $H = 0.901$ for paragliders over six cadence classes and 489 million increments, 0.885 for hang gliders, against 0.876 from the plain increment and 1.009 from the second-order variation — three routes agreeing near 0.9, and within a hundredth of the published figure. It also shows the scaling form is *not* exactly satisfied: H falls from 0.926 at the lower quartile and median to 0.872 at the upper quartile and ninetieth percentile, a spread of 0.077 against 0.004 on an exactly self-similar process — the centre of the distribution spreads faster than its flanks. Still to build: the turning-angle distribution and the speed/vertical-velocity distributions.

- **Segmentation (Sec. 3.2)**. Study Jérémie’s internship report and code, then re-examine the transition/search/climb HMM (binary features: sign of vertical speed, curvature-angle persistence, straightness) — decide whether to keep that approach or add complexity for a better segmentation on our larger, more heterogeneous data.
- **Phase-resolved observables (Sec. 3.3)**. After segmentation: the waiting-time decomposition (climb τ_c , search τ_s , and their sum τ_w), the glide step-length/duration joint law, per-flight anatomy (thermal count, phase time/distance budget), and the angular diffusion between successive glides. Plus two couplings that decide which transport model applies: the step–waiting coupling (ℓ vs τ_w , the Lévy-walk signature) and memory in the phase sequence (transition-matrix structure, duration correlations), which test the renewal assumption behind the textbook CTRW/Lévy-walk results.
- **Reproduce the reference (Sec. 3.4)**. Check that the enlarged paraglider-and-hang-glider dataset reproduces Vilpelle et al.: per-phase conditional MSD (their Fig. 3) and per-phase descriptive statistics (their Fig. 4), plus the global $H \approx 0.88$ — validation and universality test (the sailplane comparison becomes reachable once that source is in hand).
- **Minimal model (Sec. 3.5)**. A minimal stochastic model for $H \approx 0.88$: resolve transition segments, merge search+climb into one effective state, with angular diffusion between successive transition directions (preliminary idea, to be refined/changed).
- **Tooling (Sec. 3.6)**. Largely underway already: the analysis package mirrors acquisition, and the dataset statistics plus several diagnostic figures (pre-processing, altitude noise) already regenerate automatically from the data rather than being pasted in by hand. Remaining: extend the same treatment to whatever tables/figures the transport analysis produces, once it exists.